

Lab - Implement Azure Kubernetes Service

Lab scenario

- Contoso has several multi-tier applications that are not suitable to run by using Azure Container Instances.
- In order to determine whether they can be run as containerized workloads, you want to evaluate using Kubernetes as the container orchestrator.
- To further minimize management overhead, you want to test Azure Kubernetes Service, including its simplified deployment experience and scaling capabilities.

Objectives

In this lab, you will:

- Task 1: Deploy an Azure Kubernetes Service cluster
- Task 2: Deploy pods into the Azure Kubernetes Service cluster
- Task 3: Scale containerized workloads in the Azure Kubernetes service cluster

Instructions

Exercise 1

Task 1: Deploy an Azure Kubernetes Service cluster

In this task, you will deploy an Azure Kubernetes Services cluster by using the Azure portal.

1. Sign in to the [Azure portal](#).
2. In the Azure portal, search for locate **Kubernetes services** and then, on the **Kubernetes services** blade, click + **Add**.
3. On the **Basics** tab of the **Create Kubernetes cluster** blade, specify the following settings (leave others with their default values):

Setting	Value
Subscription	the name of the Azure subscription you are using in this lab
Resource group	the name of a new resource group az104-09c-rg1
Kubernetes cluster name	az104-9c-aks1
Region	the name of a region where you can provision a Kubernetes cluster
Kubernetes version	accept the default
Node size	accept the default
Node count	1

4. Click **Next: Scale >** and, on the **Scale** tab of the **Create Kubernetes cluster** blade, specify the following settings (leave others with their default values):

Setting	Value
Virtual nodes	Disabled
VM scale sets	Enabled

5. Click **Next: Authentication >** and, on the **Authentication** tab of the **Create Kubernetes cluster** blade, specify the following settings (leave others with their default values):

Setting	Value
Service principal	accept the default
Enable RBAC	Yes

6. Click **Next: Networking** > and, on the **Networking** tab of the **Create Kubernetes cluster** blade, specify the following settings (leave others with their default values):

Setting	Value
DNS name prefix	any valid, globally unique DNS host name
HTTP application routing	No
Load balancer	Standard
Network configuration	Advanced

7. Click **Next: Monitoring** >, on the **Monitoring** tab of the **Create Kubernetes cluster** blade, set **Enable container monitoring** to **No**, click **Review + create** and then click **Create**.

Note: In production scenarios, you would want to enable monitoring. Monitoring is disabled in this case since it is not covered in the lab.

Note: Wait for the deployment to complete. This should take about 10 minutes.

Task 2: Deploy pods into the Azure Kubernetes Service cluster

In this task, you will deploy a pod into the Azure Kubernetes Service cluster.

1. On the deployment blade, click the **Go to resource** link.
2. On the **az104-9c-aks1** Kubernetes service blade, in the **Settings** section, click **Node pools**.
3. On the **az104-9c-aks1 - Node pools** blade, verify that the cluster consists of a single pool with one node.

4. In the Azure portal, open the **Azure Cloud Shell** by clicking on the icon in the top right of the Azure Portal.
5. If prompted to select either **Bash** or **PowerShell**, select **Bash**.

Note: If this is the first time you are starting **Cloud Shell** and you are presented with the **You have no storage mounted** message, select the subscription you are using in this lab, and click **Create storage**.

6. From the Cloud Shell pane, run the following to retrieve the credentials to access the AKS cluster:

```
RESOURCE_GROUP='az104-09c-rg1'  
  
AKS_CLUSTER='az104-9c-aks1'  
  
az aks get-credentials --resource-group $RESOURCE_GROUP --name  
$AKS_CLUSTER
```

7. From the **Cloud Shell** pane, run the following to verify connectivity to the AKS cluster:

```
kubectl get nodes
```

8. In the **Cloud Shell** pane, review the output and verify that the one node which the cluster consists of at this point is reporting the **Ready** status.
9. From the **Cloud Shell** pane, run the following to deploy the **nginx** image from the Docker Hub:

```
kubectl create deployment nginx-deployment --image=nginx
```

Note: Make sure to use lower case letters when typing the name of the deployment (nginx-deployment)

10. From the **Cloud Shell** pane, run the following to verify that a Kubernetes pod has been created:

```
kubectl get pods
```

11. From the **Cloud Shell** pane, run the following to identify the state of the deployment:

```
kubectl get deployment
```

12. From the **Cloud Shell** pane, run the following to make the pod available from Internet:

```
kubectl expose deployment nginx-deployment --port=80 --type=LoadBalancer
```

13. From the **Cloud Shell** pane, run the following to identify whether a public IP address has been provisioned:

```
kubectl get service
```

14. Re-run the command until the value in the **EXTERNAL-IP** column for the **nginx-deployment** entry changes from **<pending>** to a public IP address. Note the public IP address in the **EXTERNAL-IP** column for **nginx-deployment**.
15. Open a browser window and navigate to the IP address you obtained in the previous step. Verify that the browser page displays the **Welcome to nginx!** message.

Task 3: Scale containerized workloads in the Azure Kubernetes service cluster

In this task, you will scale horizontally the number of pods and then number of cluster nodes.

1. From the **Cloud Shell** pane, run the following to scale the deployment by increasing of the number of pods to 2:

```
kubectl scale --replicas=2 deployment/nginx-deployment
```

2. From the **Cloud Shell** pane, run the following to verify the outcome of scaling the deployment:

```
kubectl get pods
```

Note: Review the output of the command and verify that the number of pods increased to 2.

3. From the **Cloud Shell** pane, run the following to scale out the cluster by increasing the number of nodes to 2:

```
az aks scale --resource-group $RESOURCE_GROUP --name $AKS_CLUSTER --node-count 2
```

Note: Wait for the provisioning of the additional node to complete. This might take about 3 minutes. If it fails, rerun the `az aks scale` command.

- From the **Cloud Shell** pane, run the following to verify the outcome of scaling the cluster:

```
kubectl get nodes
```

Note: Review the output of the command and verify that the number of nodes increased to 2.

- From the **Cloud Shell** pane, run the following to scale the deployment:

```
kubectl scale --replicas=10 deployment/nginx-deployment
```

- From the **Cloud Shell** pane, run the following to verify the outcome of scaling the deployment:

```
kubectl get pods
```

Note: Review the output of the command and verify that the number of pods increased to 10.

- From the **Cloud Shell** pane, run the following to review the pods distribution across cluster nodes:

```
kubectl get pod -o=custom-columns=NODE:.spec.nodeName,POD:.metadata.name
```

Note: Review the output of the command and verify that the pods are distributed across both nodes.

- From the **Cloud Shell** pane, run the following to delete the deployment:

```
kubectl delete deployment nginx-deployment
```

- Close the **Cloud Shell** pane.

Clean up resources

Note: Remember to remove any newly created Azure resources that you no longer use. Removing unused resources ensures you will not see unexpected charges.

- In the Azure portal, open the **Bash** shell session within the **Cloud Shell** pane.
- List all resource groups created throughout the labs of this module by running the following command:

ShellCopy

```
az group list --query "[?starts_with(name,'az104-09c')].name" --output tsv
```

3. Delete all resource groups you created throughout the labs of this module by running the following command:

ShellCopy

```
az group list --query "[?starts_with(name,'az104-09c')].[name]" --output  
tsv | xargs -L1 bash -c 'az group delete --name $0 --no-wait --yes'
```

Note: The command executes asynchronously (as determined by the `--no-wait` parameter), so while you will be able to run another Azure CLI command immediately afterwards within the same Bash session, it will take a few minutes before the resource groups are actually removed.

Review

In this lab, you have:

- Deployed an Azure Kubernetes Service cluster
- Deployed pods into the Azure Kubernetes Service cluster
- Scaled containerized workloads in the Azure Kubernetes service cluster